


Training dynamics: per-trajectory + rolling-mean (window=25)

Trajectory reward over training



Binary success over training

