


Training dynamics from the per-trajectory log (n=457 optimized trajectories)

Trajectory reward over training (n=457)



Binary success over training

