


Held-out evaluation: Base SFT vs RL-on-filtered (n=100, paired)

Mean judge success score (judge.success in [0, 1])



Mean composite reward

