



PPO training diagnostics


Policy loss



Entropy



KL divergence (within-trajectory)



Reference-KL (vs base SFT)

